

# Research & Engineering Insight Document

## AgentCheck: AI Agent for Automated Qualification Verification

### Table of Contents

1. Research Insight
2. Engineering Solution Thinking
3. Real-World Applicability
4. Security & Compliance Thoughts

## 1. Research Insight

### 1.1 AI Agent Patterns Studied

During the design phase, I researched several prominent AI agent architectures:

Pattern	Description	Pros	Cons
<b>ReAct (Reasoning + Acting)</b>	Agent alternates between reasoning and taking actions	Clear thought process, interpretable	Can be verbose, slower
<b>Function Calling</b>	LLM decides which tool to use via structured outputs	Reliable, deterministic tool selection	Limited reasoning visibility
<b>Plan-and-Execute</b>	Agent creates plan first, then executes steps	Good for complex tasks	Plan may become stale
<b>Multi-Agent Systems</b>	Specialized agents collaborate	Modularity, separation of concerns	Coordination overhead
<b>Tool-Use Agents</b>	Agent has access to tools and decides when to use them	Flexible, extensible	Requires good tool design

### 1.2 Why I Chose Multi-Agent + Function Calling

I selected a **Multi-Agent Architecture with Function Calling** for the following reasons:

1. **Separation of Concerns:** Each agent has a clear responsibility
  - Extraction Agent: Document processing expertise
  - Email Agent: Communication handling
  - Decision Agent: Compliance reasoning

2. **Auditability:** In RegTech, audit trails are critical. Having distinct agents makes it easy to track who did what and why.
3. **Reliability:** Function calling provides structured tool invocation, reducing hallucination risks compared to free-form agent reasoning.
4. **Modularity:** Agents can be improved, tested, or replaced independently.
5. **Compliance Alignment:** The workflow mirrors how a human compliance officer would work, making it easier to explain to auditors.

### 1.3 Alternatives Considered

#### Alternative 1: Single Monolithic Agent

```
PDF → Single Agent (does everything) → Report
```

##### Why Rejected:

- Hard to audit individual steps
- Single point of failure
- Difficult to maintain and test
- Violates single responsibility principle

#### Alternative 2: Pure ReAct Pattern

```
While not complete:  
Thought → Action → Observation → Thought → ...
```

##### Why Rejected:

- Less deterministic
- Harder to guarantee all required steps are executed
- More expensive (many LLM calls)
- Audit trail less structured

#### Alternative 3: Fully Autonomous Agent

Let the agent decide the entire workflow dynamically.

##### Why Rejected:

- Too unpredictable for compliance use cases
- Regulatory requirements need deterministic behavior
- Harder to explain decisions to auditors

## 1.4 Trade-offs Analysis

Aspect	Our Approach	Trade-off
<b>Determinism</b>	High - Fixed workflow steps	Less flexibility for edge cases
<b>Auditability</b>	Excellent - Every step logged	More storage/logging overhead
<b>Tool Execution</b>	Structured function calls	Need to maintain tool definitions
<b>LLM Dependency</b>	For extraction/analysis only	Graceful fallback when LLM unavailable
<b>Agent Autonomy</b>	Limited - Orchestrator controls flow	May miss optimization opportunities

## 2. Engineering Solution Thinking

### 2.1 Architecture Decisions

#### Decision 1: Tools as First-Class Citizens

Each tool is a discrete, testable unit:

```
def parse_pdf(pdf_path: str) -> Dict[str, Any]
def extract_fields(raw_text: str) -> ExtractedFields
def identify_university(extracted_fields: ExtractedFields) -> str
def lookup_contact(university_name: str) -> Optional[UniversityContact]
def draft_email(...) -> Dict[str, str]
def send_to_outbox(...) -> OutgoingEmail
def read_reply(...) -> IncomingEmail
def analyze_reply(...) -> ReplyAnalysis
def decide_compliance(...) -> Tuple[ComplianceResult, str]
```

#### Rationale:

- Tools can be unit tested independently
- Easy to add new tools without modifying agents
- Clear interface contracts
- Audit logger wraps each tool call automatically

#### Decision 2: Prompt Templates (Jinja2)

Prompts are externalized to `config/prompts/*.j2`:

```
# extract_fields.j2
Extract the following fields from the certificate text:
- candidate_name: {{ instructions }}
...
Certificate Text: {{ certificate_text }}
```

### Rationale:

- Non-engineers can modify prompts
- Version control for prompt changes
- Easy A/B testing of prompts
- Separation of logic and content

### Decision 3: Config-Driven University Mappings

```
{  
  "universities": {  
    "University of Example": {  
      "email": "verification@example.edu",  
      "country": "USA",  
      "verification_department": "Office of the Registrar"  
    }  
  }  
}
```

### Rationale:

- Business users can add universities without code changes
- Easy to integrate with external data sources
- Supports different verification endpoints per university

## 2.2 Balancing Hard-Coded Logic vs. Generative Reasoning

Step	Approach	Reason
Workflow order	Hard-coded	Must be consistent for audit
Field extraction	LLM	Handles variations in certificate formats
University lookup	Database + fuzzy match	Deterministic, fast
Email drafting	LLM	Natural language generation
Reply analysis	LLM	Handles varied response formats
Compliance mapping	Hard-coded	Regulatory requirement

**Key Insight:** Use LLM where natural language understanding is needed; use deterministic code where consistency is required.

## 2.3 Failure Cases and Mitigations

Failure Case	Mitigation
PDF unreadable	Fallback to text file; return INCONCLUSIVE

Failure Case	Mitigation
LLM unavailable	Mock responses for demo; queue for retry
University not found	Mark as INCONCLUSIVE; log for manual review
Ambiguous reply	Lower confidence score; may need human review
LLM hallucination	Use structured output; validate response schema
Network timeout	Retry with exponential backoff
Rate limiting	Task queue with throttling

## 2.4 Audit Trail Design

Every action creates an `AuditLogEntry`:

```
@dataclass
class AuditLogEntry:
    timestamp: datetime
    step: str          # "001_parse_pdf"
    action: str        # Human-readable description
    agent: str         # Which agent performed this
    tool: str          # Which tool was used
    input_data: Dict   # Sanitized inputs
    output_data: Dict # Sanitized outputs
    success: bool
    error_message: Optional[str]
```

### Key Features:

- Sequential step numbering for timeline reconstruction
- Sensitive data automatically redacted
- Stored as JSONL for streaming writes
- Summary file generated at session end

## 3. Real-World Applicability

### 3.1 Extension to Real Email Integration

**Current:** Simulated outbox/inbox

#### Production Path:

```
class RealEmailService(EmailService):
    def send_email(self, email: OutgoingEmail) -> str:
        # Use SendGrid/AWS SES/SMTP
        response = sendgrid.send()
```

```

        to=email.recipient_email,
        subject=email.subject,
        body=email.body
    )
    return response.message_id

def poll_inbox(self, reference_id: str) -> Optional[IncomingEmail]:
    # Use IMAP/Gmail API/Webhook
    messages = imap.search(subject=f"RE: {reference_id}")
    return self._parse_message(messages[0]) if messages else None

```

## Considerations:

- Email sending rate limits
- Reply detection (subject matching, thread tracking)
- Spam folder handling
- Bounce handling

## 3.2 Real Certificate OCR

**Current:** Text extraction from PDF

### Production Enhancement:

```

class OCEReaderParser(PDFParser):
    def parse_pdf(self, pdf_path: str) -> Dict:
        # Try text extraction first
        text = self._extract_text(pdf_path)

        if len(text.strip()) < 50: # Likely scanned image
            text = self._ocr_extract(pdf_path)

        return {"raw_text": text, "ocr_used": True}

    def _ocr_extract(self, pdf_path: str) -> str:
        # Use Tesseract/AWS Textract/Google Vision
        return textract.process(pdf_path)

```

## 3.3 API-Based University Verification

Many universities offer verification APIs:

```

class UniversityAPIVerifier:
    SUPPORTED_APIS = {
        "national_clearinghouse": NationalClearinghouseAPI,
        "hedd_uk": HEDDVerificationAPI,
        "direct_api": DirectUniversityAPI
    }

```


```

async def verify(self, certificate: ExtractedFields) -> VerificationResult:
    api = self._get_api_for_university(certificate.university_name)
    return await api.verify(
        name=certificate.candidate_name,
        degree=certificate.degree_name,
        date=certificate.issue_date
    )

```

### 3.4 Scaling to 1,000+ Checks/Day

#### Architecture for Scale:



#### Key Components:

- **Redis Queue:** Decouple request from processing
- **Celery Workers:** Horizontal scaling
- **Rate Limiter:** Per-university rate limiting
- **Retry Logic:** Exponential backoff
- **Dead Letter Queue:** Failed tasks for manual review

### 3.5 Monitoring, SLA, and Retries

```

# Example monitoring setup
class VerificationMetrics:
    def __init__(self):
        self.processing_time = Histogram('verification_duration_seconds')
        self.result_counter = Counter('verification_results_total', ['result'])
        self.error_counter = Counter('verification_errors_total',
                                     ['error_type'])

    def track_verification(self, report: ComplianceReport):
        self.processing_time.observe(report.processing_time_seconds)
        self.result_counter.labels(result=report.compliance_result.value).inc()

# SLA Configuration
SLA_CONFIG = {

```

```

    "max_processing_time_seconds": 30,
    "max_retry_attempts": 3,
    "retry_backoff_seconds": [10, 30, 60],
    "alert_threshold_queue_size": 100
}

```

## 4. Security & Compliance Thoughts

### 4.1 Data Privacy

**Personal Data Handled:**

- Candidate names
- Degree information
- Certificate images

**Protections Implemented:**

- Audit logs sanitize sensitive data
- No PII in log messages
- Data stored locally (no third-party storage)

**Production Recommendations:**

- Encrypt data at rest (AES-256)
- Encrypt in transit (TLS 1.3)
- Implement data retention policies
- GDPR/CCPA compliance (right to deletion)
- PII masking in logs

### 4.2 Model Hallucination Risks

**Identified Risks:**

1. LLM fabricates certificate fields
2. LLM misinterprets university reply
3. LLM invents confidence scores

**Mitigations:**

```

# 1. Schema validation
def extract_fields(raw_text: str) -> ExtractedFields:
    response = llm.complete_json(prompt)
    # Validate against known patterns
    if not self._validate_date_format(response.get("issue_date")):
        response["issue_date"] = None
    return ExtractedFields(**response)

```

```

# 2. Confidence thresholds
if analysis.confidence_score < 0.7:
    result = VerificationStatus.INCONCLUSIVE
    # Flag for human review

# 3. Fallback to keyword analysis
def _fallback_analyze_reply(self, reply_text: str) -> ReplyAnalysis:
    # Deterministic keyword matching as backup
    ...

```

## 4.3 Traceability Requirements

### Compliance Standard Alignment:

Requirement	Implementation
Who made the decision?	agent field in audit log
When?	timestamp field (UTC)
What inputs?	input_data field
What outputs?	output_data field
Why?	decision_explanation field
Can it be reproduced?	Session ID links all related logs

### Report Retention:

- Reports stored as JSON files
- Session summaries for quick lookup
- JSONL logs for detailed reconstruction

## 4.4 Human-in-the-Loop Necessity

### When Human Review is Required:

#### 1. Low Confidence Results

```

if analysis.confidence_score < CONFIDENCE_THRESHOLD:
    report.requires_human_review = True

```

#### 2. INCONCLUSIVE Results

- University not in database
- Ambiguous reply received
- PDF parsing failed


#### 3. High-Stakes Decisions

- NOT\_COMPLIANT results affecting employment
- First-time university verification

#### 4. Anomaly Detection

```
if self._detect_anomaly(certificate):
    # Unusual patterns trigger review
    report.flagged_for_review = True
```

#### Production Workflow:



## Conclusion

This document outlines the research, engineering decisions, and practical considerations for the AgentCheck certificate verification system. The multi-agent architecture with function calling provides a balance of:

- **Reliability** through structured tool execution
- **Flexibility** through LLM-powered analysis
- **Auditability** through comprehensive logging
- **Extensibility** through modular design

For production deployment, the key enhancements would be:

1. Real email integration
2. University API connections
3. Enhanced OCR capabilities
4. Horizontal scaling with message queues
5. Robust human-in-the-loop workflows

The system is designed with RegTech compliance requirements at its core, ensuring that every decision can be traced, explained, and audited.

---

*Document Version: 1.0.0*

*Last Updated: December 2024*