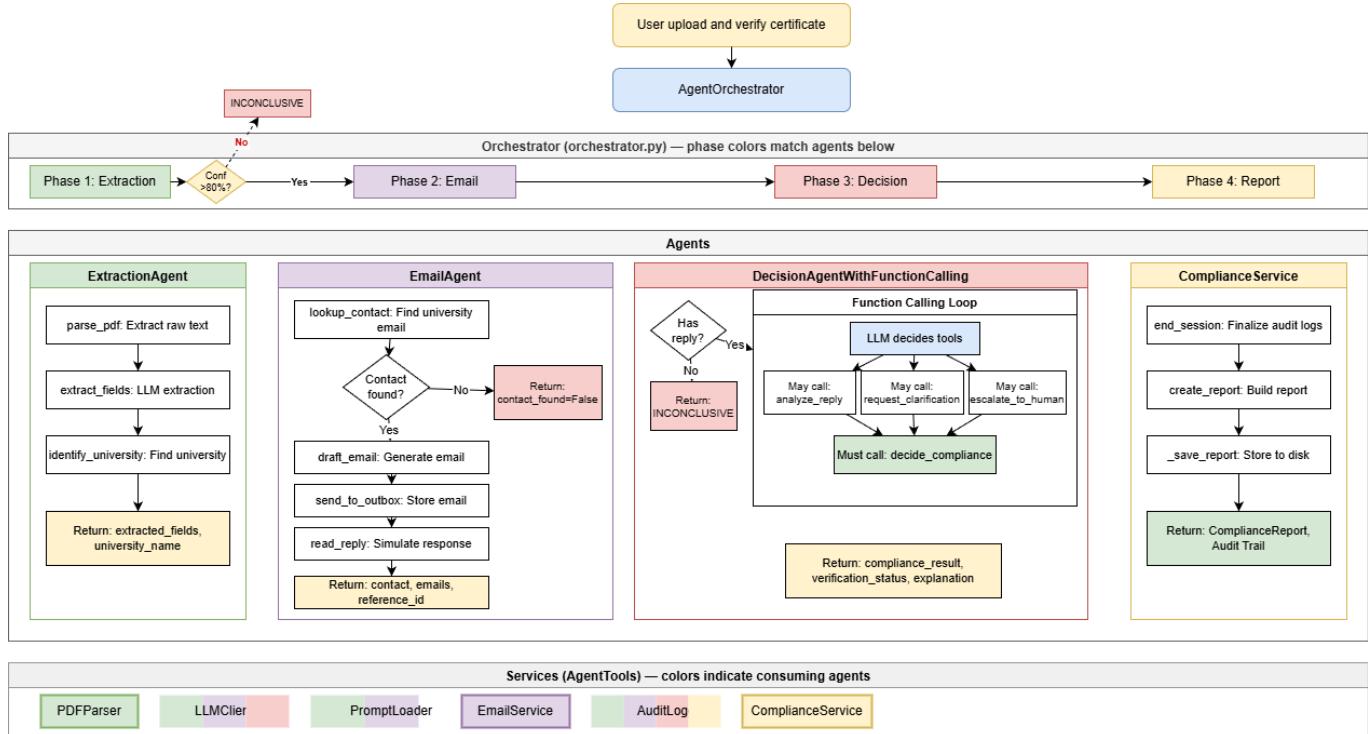


Research & Engineering Insight

1. Research Insight

Architecture Overview

The workflow follows a fixed sequential order: **Extraction** → **Email** → **Decision** → **Report**. Compliance processes require predictable, auditable execution paths—the Orchestrator enforces this strictly.



Three specialized agents handle each domain:

| Agent | LLM Usage | Behavior |
|------------------------|---------------------------------------|------------------------------------|
| ExtractionAgent | Content generation (field extraction) | Deterministic tool sequence |
| EmailAgent | Content generation (email drafting) | Deterministic tool sequence |
| DecisionAgent | Tool selection via Function Calling | Adaptive—handles ambiguous replies |

This separation enables clear audit trails—each agent's actions are logged independently with timestamps and inputs/outputs.

Why Handcrafted Orchestration + Function Calling?

The alternative was full LangChain/LangGraph—letting the framework handle both orchestration and tool execution. Given the prototype scope and timeline, the **handcrafted approach** was chosen: zero additional dependencies, and the LLM decides only where needed instead of every step.

LangChain excels at dynamic agent behavior—and with more time, it could be configured to enforce sequential phases. However, for a fixed compliance workflow, handcrafted orchestration achieves the same result with less complexity. **Function Calling in DecisionAgent** provides the balance: the workflow stays deterministic, but the LLM retains intelligent tool selection where ambiguity exists (interpreting email replies). This avoids both extremes—fully autonomous agents and fully hardcoded logic. If requirements evolve (e.g., dynamic phase ordering), LangGraph would be reconsidered.

Trade-offs Accepted

| What I Gave Up | Why It's Acceptable |
|-----------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------|
| Dynamic workflow changes—cannot reorder phases at runtime | Compliance workflows are inherently fixed; dynamic ordering would break audit requirements |
| Framework ecosystem—no LangChain integrations (memory, callbacks, tracing) | Prototype scope doesn't require these; can add later if needed |
| Less "agent-like"—most steps are deterministic, not LLM-driven | Deliberate choice: predictability over autonomy for regulated processes |
| Function Calling in DecisionAgent compensates for the third point—the system retains intelligent decision-making where it matters most. | |

2. Engineering Solution Thinking

Tools, Prompts, and Models Structure

Tools use a mixin pattern to group domain-specific logic, simplifying code structure. **Prompts** are externalized to **Jinja2 templates** (`config/prompts/*.j2`) for independent versioning. **University Mappings** are config-driven via `config/universities.json`—no code deployment needed for updates.

Failure Cases

| Failure | Handling |
|-------------------------|--------------------------------------------------------------------------------------------------|
| PDF unreadable | LLM Vision API extracts text from rendered images; handles scanned documents |
| Low-quality/damaged PDF | Vision API detects visual damage; extraction confidence $\leq 80\%$ triggers INCONCLUSIVE |
| University not found | Return INCONCLUSIVE ; flag for manual review |
| Ambiguous/fraud reply | Lower confidence score triggers human review queue |
| LLM hallucination | Schema validation on all outputs; reject malformed responses |

The decision to use **LLM Vision for all PDFs** (instead of traditional Python PDF readers like PyMuPDF) was deliberate: it avoids arbitrary fallback thresholds and provides consistent handling regardless of PDF type—digital or scanned.

Auditability

Every tool call creates an **AuditLogEntry** with: timestamp, agent name, tool name, sanitized inputs/outputs, and success status. Credentials are redacted and long values truncated in logs. PII sanitization is extensible via config. The audit trail enables **timeline reconstruction** (sequential step numbering), **accountability** (which agent made which decision), and **reproducibility** (Session ID links all related logs).

3. Real-World Applicability

Production Extensions

| Current (Prototype) | Production Path |
|------------------------|--------------------------------------------------------------------|
| Simulated outbox/inbox | SendGrid/AWS SES for sending; IMAP/Gmail API for receiving |
| LLM Vision for PDFs | Already production-ready; handles scanned certificates |
| JSON university config | Database + external APIs (My eQuals, National Clearinghouse, etc.) |

Scaling to 1,000+ Checks/Day

Architecture: FastAPI → Redis Queue → Celery Workers. Features include horizontal scaling via worker pool, per-university rate limiting (avoid spam triggers), and dead letter queue for failed verifications.

Monitoring & Reliability: Target 30s processing time (SLA); alert if queue exceeds threshold. Retries use exponential backoff (10s → 30s → 60s) for failures. Observability via Prometheus metrics on processing time, result distribution, and error rates.

4. Security & Compliance Thoughts

Data Privacy

Audit logs automatically redact sensitive credentials. Production plans include standard encryption (AES-256, TLS 1.3) and GDPR/CCPA-compliant retention policies.

Model Hallucination Risks

| Mitigation | Implementation |
|------------------------------|-----------------------------------------------------------------------------------------------|
| Schema Validation | Pydantic models validate all LLM outputs to ensure structural integrity |
| Structured Outputs | Function Calling restricts responses to valid enum values, preventing open-ended fabrications |
| Extraction Confidence | Vision API assesses document quality; confidence ≤80% escalates to human review |
| Decision Confidence | Reply analysis scores below 0.7 trigger fallback to human review |

Traceability

The detailed audit trail allows for complete timeline reconstruction. A unique **Session ID** links every log entry for a specific verification request, ensuring that every automated decision can be traced back to its inputs and timestamp.

Human-in-the-Loop

This architecture prioritizes safety by treating the AI as an **accelerator, not a replacement**. Full automation is strictly reserved for clear-cut success paths.

Automatic Breaks are triggered at two checkpoints:

1. **Extraction Phase:** Document quality confidence ≤80% immediately returns **INCONCLUSIVE**—damaged or altered documents are flagged before verification proceeds.
2. **Decision Phase:** Any result marked **INCONCLUSIVE / NOT_COMPLIANT**, or **COMPLIANT** with low confidence, is routed to a human review queue.

Only high-confidence extractions from high-quality documents with clear **COMPLIANT** responses bypass this safety layer.