

Research & Engineering Insight Document

AgentCheck: AI Agent for Automated Qualification Verification

Table of Contents

1. Design Considerations & Evolution
 2. Research Insight
 3. Engineering Solution Thinking
 4. Real-World Applicability
 5. Security & Compliance Thoughts
-

1. Design Considerations & Evolution

1.1 The Framework Question

The assignment offers multiple approaches: LangChain, Llamaindex, OpenAI function calling, or a handcrafted agent loop. My first instinct was to reach for LangChain—it's popular, well-documented, and would demonstrate familiarity with modern tooling.

But I paused to ask: **What does this project actually need?**

The scope is clear: 3 agents, ~10 tools, 1 linear workflow, prototype scale. LangChain brings 50+ dependencies, complex abstractions, and patterns designed for enterprise-scale agent orchestration. Using it here would be like bringing a forklift to move a chair—technically capable, but not the right tool.

More importantly, in an interview context, wrapping everything in LangChain would hide whether I actually understand how agents work. Anyone can call `AgentExecutor.invoke()`. Fewer can explain the reasoning loop underneath.

1.2 My Decision: Handcrafted + Targeted Enhancement

I chose a **hybrid approach**:

Component	Approach	Rationale
Overall workflow	Fixed pipeline	Compliance needs predictability
Extraction & Email agents	Deterministic steps	These tasks are mechanical, not ambiguous
Decision agent	OpenAI Function Calling	This is where flexibility actually matters

Why this split?

Parsing a PDF is parsing a PDF—there's no benefit to letting an LLM decide whether to parse it. But interpreting a university's reply? That's genuinely ambiguous. A reply might be clear confirmation, a request for more documents, or something suspicious that needs human eyes.

The Decision agent is the only place where dynamic tool selection provides real value. Elsewhere, it would add complexity without benefit.

1.3 Why Function Calling Over LangChain

Consideration	Function Calling	LangChain
Dependencies	Zero new packages	50+ packages
Debugging	Direct stack traces	Framework abstractions
Interview signal	Shows understanding	Shows usage
Right-sized?	Yes—matches scope	Over-engineered
Future flexibility	Can add LangChain later	Already committed

The key insight: **I can always add LangChain if scale demands it. I can't easily remove it once embedded.**

1.4 What the Enhancement Enables

With Function Calling, the Decision agent gains four capabilities:

1. **Analyze** — Understand what the university is communicating
2. **Clarify** — Flag when more information is needed
3. **Escalate** — Route suspicious cases to humans
4. **Decide** — Make the final compliance call

The LLM chooses which to invoke based on context. A clear "certificate confirmed" reply goes straight to decision. An ambiguous reply triggers analysis first. A suspicious sender domain triggers escalation—no auto-decision at all.

This isn't just technical flexibility. It's **appropriate caution** for a compliance system.

1.5 Trade-offs Acknowledged

Nothing is free. This approach means:

- **More LLM calls** for complex cases (cost/latency)
- **Less predictable** behavior than pure fixed pipeline
- **More logging** needed to maintain audit trail

I accepted these trade-offs because:

- Complex cases are rare—most are straightforward
- Unpredictability is bounded (max 5 iterations, terminal conditions)
- Logging is already built in—marginal cost to extend

1.6 Evolution Path

This architecture is designed to grow:

Now:	Handcrafted + Function Calling (right for prototype)
If needed:	Add LangChain for multi-workflow orchestration
Later:	Full agent framework with memory and retrieval

The principle: **Start simple, add complexity when earned.**

2. Research Insight

2.1 AI Agent Patterns Studied

During the design phase, I researched several prominent AI agent architectures:

Pattern	Description	Pros	Cons
ReAct (Reasoning + Acting)	Agent alternates between reasoning and taking actions	Clear thought process, interpretable	Can be verbose, slower
Function Calling	LLM decides which tool to use via structured outputs	Reliable, deterministic tool selection	Limited reasoning visibility
Plan-and-Execute	Agent creates plan first, then executes steps	Good for complex tasks	Plan may become stale
Multi-Agent Systems	Specialized agents collaborate	Modularity, separation of concerns	Coordination overhead
Tool-Use Agents	Agent has access to tools and decides when to use them	Flexible, extensible	Requires good tool design

2.2 Why I Chose Multi-Agent + Function Calling

I selected a **Multi-Agent Architecture with Function Calling** for the following reasons:

1. **Separation of Concerns:** Each agent has a clear responsibility

- Extraction Agent: Document processing expertise
- Email Agent: Communication handling
- Decision Agent: Compliance reasoning

2. **Auditability:** In RegTech, audit trails are critical. Having distinct agents makes it easy to track who did what and why.

3. **Reliability:** Function calling provides structured tool invocation, reducing hallucination risks compared to free-form agent reasoning.

4. **Modularity:** Agents can be improved, tested, or replaced independently.

- 5. Compliance Alignment:** The workflow mirrors how a human compliance officer would work, making it easier to explain to auditors.

2.3 Alternatives Considered

Alternative 1: Single Monolithic Agent

```
PDF → Single Agent (does everything) → Report
```

Why Rejected:

- Hard to audit individual steps
- Single point of failure
- Difficult to maintain and test
- Violates single responsibility principle

Alternative 2: Pure ReAct Pattern

```
While not complete:  
Thought → Action → Observation → Thought → ...
```

Why Rejected:

- Less deterministic
- Harder to guarantee all required steps are executed
- More expensive (many LLM calls)
- Audit trail less structured

Alternative 3: Fully Autonomous Agent

Let the agent decide the entire workflow dynamically.

Why Rejected:

- Too unpredictable for compliance use cases
- Regulatory requirements need deterministic behavior
- Harder to explain decisions to auditors

2.4 Trade-offs Analysis

Aspect	Our Approach	Trade-off
Determinism	High - Fixed workflow steps	Less flexibility for edge cases
Auditability	Excellent - Every step logged	More storage/logging overhead

Aspect	Our Approach	Trade-off
Tool Execution	Structured function calls	Need to maintain tool definitions
LLM Dependency	For extraction/analysis only	Graceful fallback when LLM unavailable
Agent Autonomy	Limited - Orchestrator controls flow	May miss optimization opportunities

3. Engineering Solution Thinking

3.1 Architecture Decisions

Decision 1: Tools as First-Class Citizens

Each tool is a discrete, testable unit:

```
def parse_pdf(pdf_path: str) -> Dict[str, Any]
def extract_fields(raw_text: str) -> ExtractedFields
def identify_university(extracted_fields: ExtractedFields) -> str
def lookup_contact(university_name: str) -> Optional[UniversityContact]
def draft_email(...) -> Dict[str, str]
def send_to_outbox(...) -> OutgoingEmail
def read_reply(...) -> IncomingEmail
def analyze_reply(...) -> ReplyAnalysis
def decide_compliance(...) -> Tuple[ComplianceResult, str]
```

Rationale:

- Tools can be unit tested independently
- Easy to add new tools without modifying agents
- Clear interface contracts
- Audit logger wraps each tool call automatically

Decision 2: Prompt Templates (Jinja2)

Prompts are externalized to `config/prompts/*.j2`:

```
# extract_fields.j2
Extract the following fields from the certificate text:
- candidate_name: {{ instructions }}
...
Certificate Text: {{ certificate_text }}
```

Rationale:

- Non-engineers can modify prompts
- Version control for prompt changes

- Easy A/B testing of prompts
- Separation of logic and content

Decision 3: Config-Driven University Mappings

```
{
  "universities": {
    "University of Example": {
      "email": "verification@example.edu",
      "country": "USA",
      "verification_department": "Office of the Registrar"
    }
  }
}
```

Rationale:

- Business users can add universities without code changes
- Easy to integrate with external data sources
- Supports different verification endpoints per university

3.2 Balancing Hard-Coded Logic vs. Generative Reasoning

Step	Approach	Reason
Workflow order	Hard-coded	Must be consistent for audit
Field extraction	LLM	Handles variations in certificate formats
University lookup	Database + fuzzy match	Deterministic, fast
Email drafting	LLM	Natural language generation
Reply analysis	LLM	Handles varied response formats
Compliance mapping	Hard-coded	Regulatory requirement

Key Insight: Use LLM where natural language understanding is needed; use deterministic code where consistency is required.

3.3 Failure Cases and Mitigations

Failure Case	Mitigation
PDF unreadable	LLM Vision API extracts text from images; clear error message
LLM unavailable	Queue for retry; require API key configuration
University not found	Mark as INCONCLUSIVE; log for manual review
Ambiguous reply	Lower confidence score; may need human review

Failure Case	Mitigation
LLM hallucination	Use structured output; validate response schema
Network timeout	Retry with exponential backoff
Rate limiting	Task queue with throttling

3.4 Audit Trail Design

Every action creates an `AuditLogEntry`:

```
@dataclass
class AuditLogEntry:
    timestamp: datetime
    step: str          # "001_parse_pdf"
    action: str        # Human-readable description
    agent: str         # Which agent performed this
    tool: str          # Which tool was used
    input_data: Dict   # Sanitized inputs
    output_data: Dict # Sanitized outputs
    success: bool
    error_message: Optional[str]
```

Key Features:

- Sequential step numbering for timeline reconstruction
- Sensitive data automatically redacted
- Stored as JSONL for streaming writes
- Summary file generated at session end

4. Real-World Applicability

4.1 Extension to Real Email Integration

Current: Simulated outbox/inbox

Production Path:

```
class RealEmailService(EmailService):
    def send_email(self, email: OutgoingEmail) -> str:
        # Use SendGrid/AWS SES/SMTP
        response = sendgrid.send(
            to=email.recipient_email,
            subject=email.subject,
            body=email.body
        )
        return response.message_id
```

```

def poll_inbox(self, reference_id: str) -> Optional[IncomingEmail]:
    # Use IMAP/Gmail API/Webhook
    messages = imap.search(subject=f"RE: {reference_id}")
    return self._parse_message(messages[0]) if messages else None

```

Considerations:

- Email sending rate limits
- Reply detection (subject matching, thread tracking)
- Spam folder handling
- Bounce handling

4.2 PDF Text Extraction - Design Decision

Design Choice: Always use LLM Vision API for certificate text extraction.

Alternatives Considered:

Approach	Pros	Cons
PyMuPDF text extraction	Fast, free, no API cost	Only works for digital PDFs; fails on scanned documents
PyMuPDF + OCR fallback	Works for both types	Complex logic; arbitrary threshold for "when to fallback"
LLM Vision only <input checked="" type="checkbox"/>	Consistent, accurate, handles all PDF types	Uses API tokens

Why We Chose LLM Vision:

1. **Correctness over cost** - For compliance verification, accuracy is critical. Traditional Python text extraction (PyMuPDF) works well for digitally-created PDFs but fails completely on scanned certificates—which most real-world certificates are.
2. **No arbitrary thresholds** - Fallback approaches require deciding "when is extracted text insufficient?" (e.g., < 50 characters). This is fragile and can misclassify legitimate short certificates.
3. **Simpler implementation** - One consistent path for all PDFs is easier to maintain, test, and debug than conditional fallback logic.
4. **Better for certificates** - LLM Vision understands document layout, can read watermarks, seals, and handwritten elements that traditional extractors miss.

Implementation:

```

# PDF → Image (PyMuPDF renders at 2x zoom)
# Image → LLM Vision API → Extracted text
def parse_pdf(self, pdf_path: str) -> dict:

```

```

doc = fitz.open(pdf_path)
for page in doc:
    image = page.get_pixmap(matrix=fitz.Matrix(2, 2))
    text = self.llm_client.extract_text_from_image(image)
return {"raw_text": text, "method": "vision_api"}

```

4.3 API-Based University Verification

Many universities offer verification APIs:

```

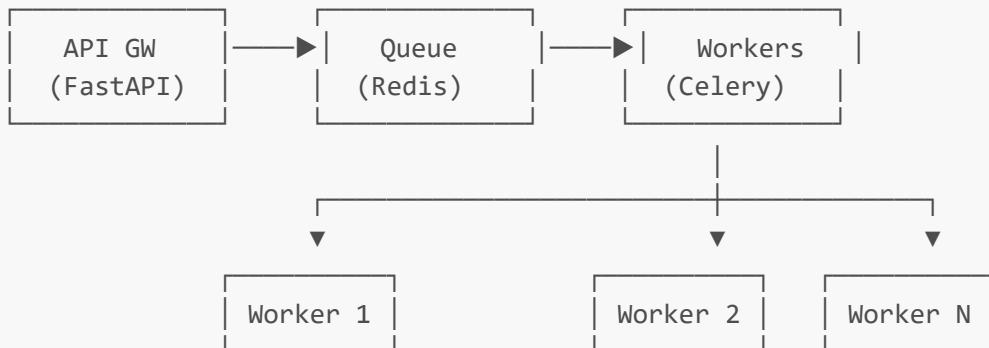
class UniversityAPIVerifier:
    SUPPORTED_APIS = {
        "national_clearinghouse": NationalClearinghouseAPI,
        "hedd_uk": HEDDVerificationAPI,
        "direct_api": DirectUniversityAPI
    }

    async def verify(self, certificate: ExtractedFields) -> VerificationResult:
        api = self._get_api_for_university(certificate.university_name)
        return await api.verify(
            name=certificate.candidate_name,
            degree=certificate.degree_name,
            date=certificate.issue_date
        )

```

4.4 Scaling to 1,000+ Checks/Day

Architecture for Scale:



Key Components:

- **Redis Queue:** Decouple request from processing
- **Celery Workers:** Horizontal scaling
- **Rate Limiter:** Per-university rate limiting
- **Retry Logic:** Exponential backoff

- **Dead Letter Queue:** Failed tasks for manual review

4.5 Monitoring, SLA, and Retries

```
# Example monitoring setup
class VerificationMetrics:
    def __init__(self):
        self.processing_time = Histogram('verification_duration_seconds')
        self.result_counter = Counter('verification_results_total', ['result'])
        self.error_counter = Counter('verification_errors_total',
                                     ['error_type'])

    def track_verification(self, report: ComplianceReport):
        self.processing_time.observe(report.processing_time_seconds)
        self.result_counter.labels(result=report.compliance_result.value).inc()

# SLA Configuration
SLA_CONFIG = {
    "max_processing_time_seconds": 30,
    "max_retry_attempts": 3,
    "retry_backoff_seconds": [10, 30, 60],
    "alert_threshold_queue_size": 100
}
```

5. Security & Compliance Thoughts

5.1 Data Privacy

Personal Data Handled:

- Candidate names
- Degree information
- Certificate images

Protections Implemented:

- Audit logs sanitize sensitive data
- No PII in log messages
- Data stored locally (no third-party storage)

Production Recommendations:

- Encrypt data at rest (AES-256)
- Encrypt in transit (TLS 1.3)
- Implement data retention policies
- GDPR/CCPA compliance (right to deletion)
- PII masking in logs

5.2 Model Hallucination Risks

Identified Risks:

1. LLM fabricates certificate fields
2. LLM misinterprets university reply
3. LLM invents confidence scores

Mitigations:

```
# 1. Schema validation
def extract_fields(raw_text: str) -> ExtractedFields:
    response = llm.complete_json(prompt)
    # Validate against known patterns
    if not self._validate_date_format(response.get("issue_date")):
        response["issue_date"] = None
    return ExtractedFields(**response)

# 2. Confidence thresholds
if analysis.confidence_score < 0.7:
    result = VerificationStatus.INCONCLUSIVE
    # Flag for human review

# 3. Fallback to keyword analysis
def _fallback_analyze_reply(self, reply_text: str) -> ReplyAnalysis:
    # Deterministic keyword matching as backup
    ...
```

5.3 Traceability Requirements

Compliance Standard Alignment:

Requirement	Implementation
Who made the decision?	agent field in audit log
When?	timestamp field (UTC)
What inputs?	input_data field
What outputs?	output_data field
Why?	decision_explanation field
Can it be reproduced?	Session ID links all related logs

Report Retention:

- Reports stored as JSON files
- Session summaries for quick lookup
- JSONL logs for detailed reconstruction

5.4 Human-in-the-Loop Necessity

When Human Review is Required:

1. Low Confidence Results

```
if analysis.confidence_score < CONFIDENCE_THRESHOLD:  
    report.requires_human_review = True
```

2. INCONCLUSIVE Results

- University not in database
- Ambiguous reply received
- PDF parsing failed


3. High-Stakes Decisions

- NOT_COMPLIANT results affecting employment
- First-time university verification

4. Anomaly Detection

```
if self._detect_anomaly(certificate):  
    # Unusual patterns trigger review  
    report.flagged_for_review = True
```

Production Workflow:



Conclusion

This document outlines the research, engineering decisions, and practical considerations for the AgentCheck certificate verification system. The multi-agent architecture with function calling provides a balance of:

- **Reliability** through structured tool execution
- **Flexibility** through LLM-powered analysis
- **Auditability** through comprehensive logging
- **Extensibility** through modular design

For production deployment, the key enhancements would be:

1. Real email integration
2. University API connections
3. Horizontal scaling with message queues
4. Robust human-in-the-loop workflows
5. Function Calling enhancement for DecisionAgent (see [IMPLEMENTATION_PLAN.md](#))

The system is designed with RegTech compliance requirements at its core, ensuring that every decision can be traced, explained, and audited.

Document Version: 1.1.0

Last Updated: December 2024